

# Programming Techniques for Scientific Simulations I

## Week 7: Algorithms, Data Structures, and Plotting

Detailed Lecture Notes

October 30, 2025

## Contents

<b>I Algorithms and Data Structures in C++ (week07a)</b>	<b>3</b>
<b>1 Complexity Analysis: Measuring Efficiency</b>	<b>3</b>
1.1 The Core Question . . . . .	3
1.2 Asymptotic Notation: The Language of Complexity . . . . .	4
1.2.1 O / Big-Oh (Worst Case Upper Bound) . . . . .	4
1.2.2 Ω / Big-Omega (Best Case Lower Bound) . . . . .	4
1.2.3 Θ / Big-Theta (Tight Bound) . . . . .	4
1.3 The Real-World Impact of Complexity . . . . .	5
1.4 Complexity Examples in C++ . . . . .	6
1.4.1 Example 1: Single Loop (Slide 7) . . . . .	6
1.4.2 Example 2: Nested Loop (Slide 8) . . . . .	6
1.4.3 Example 3: Search with 'break' (Slide 9) . . . . .	7
1.5 Amortized Analysis: The "Clever Way" to Grow Arrays . . . . .	7
<b>2 The C++ Standard Template Library (STL)</b>	<b>8</b>
2.1 What is the STL? . . . . .	8
2.2 Simple Utilities . . . . .	9
<b>3 STL Container Deep Dive</b>	<b>9</b>
3.1 Common Container Operations . . . . .	10
3.2 Sequence Containers . . . . .	10
3.2.1 'std::vector' — The "Smart Array" . . . . .	10
3.2.2 std::deque — The "Double-Ended Queue" . . . . .	12
3.2.3 std::list — The "Linked List" . . . . .	12
3.3 Container Adapters . . . . .	13
3.3.1 std::stack . . . . .	13
3.3.2 std::queue . . . . .	14

3.3.3	<code>std::priority_queue</code>	14
3.4	Associative Containers (Trees)	14
3.4.1	<code>std::set</code>	15
3.4.2	<code>std::map</code>	15
<b>4</b>	<b>The Magic Glue: Iterators</b>	<b>16</b>
4.1	The $N \times M$ Problem	16
4.2	The Solution: Generic Traversal	16
4.3	How Iterators are Implemented	17
4.4	Complete Iterator Operations Reference	18
4.4.1	Essential Iterator Operators	18
4.4.2	Complete Example: Using All Iterator Operations	19
4.5	Iterator Categories	20
<b>5</b>	<b>The Generic Algorithms</b>	<b>22</b>
5.1	Example: <code>std::find</code>	22
5.2	Example: <code>std::find_if</code>	22
5.3	Tricky Case 1: Member Functions as Predicates	23
5.4	Tricky Case 2: Copying into Empty Containers	23
5.5	Algorithm Naming Conventions	24
<b>6</b>	<b>Application and Summary</b>	<b>24</b>
6.1	Application: The Penna Model	24
6.2	Summary	25
<b>II</b>	<b>Plotting / Scientific Visualization (week07b)</b>	<b>26</b>
<b>7</b>	<b>Plotting Our Data</b>	<b>26</b>
7.1	Gnuplot	26
7.2	Python + Matplotlib	26
7.3	Best Practice: Python Virtual Environments	26

## Part I

# Algorithms and Data Structures in C++ (week07a)

## 1 Complexity Analysis: Measuring Efficiency

### 1.1 The Core Question

(Slide 2)

When we write a program, it's not enough for it to be *correct*. It must also be *efficient*. "Efficiency" can mean many things (memory usage, disk I/O), but most often it means **time**. We want our algorithms to be fast.

The core question of **complexity analysis** is:

*"How does the time needed for an algorithm scale with the problem size N?"*

Here,  $N$  is the "problem size." If you're sorting a list,  $N$  is the number of items. If you're searching a database,  $N$  is the number of entries. We don't measure time in seconds, because that changes when you buy a faster computer. We measure time in the *number of operations* as a *function* of  $N$ .

**Real-World Analogy: A Recipe** Imagine  $N$  is the number of guests coming to dinner.

- A  $O(1)$  (constant) algorithm: The recipe takes 30 minutes, whether you have 1 guest or 100. (e.g., "Pre-heat the oven").
- A  $O(N)$  (linear) algorithm: The recipe takes 5 minutes per guest. (e.g., "Chop 1 potato per guest"). Double the guests, double the time.
- A  $O(N^2)$  (quadratic) algorithm: You must introduce every guest to every other guest. Double the guests, and you quadruple the introduction time.

We analyze this scaling in a few different ways:

- **Worst Case Analysis:** This is the most common. It answers, "What is the *maximum* possible time your algorithm will take for a given  $N$ ?" This is a guarantee. (Analogy: "How long to find a name in the phone book?" Worst case: It's the very last name.)
- **Best Case Analysis:** What is the *minimum* possible time? (Analogy: The name is the very first one.)
- **Average Case Analysis:** What is the *typical* time? This is often more useful but is much harder to calculate.
- **Amortized Analysis:** What is the average time over a *sequence* of many operations? We will see this is very important for 'std::vector'.

## 1.2 Asymptotic Notation: The Language of Complexity

(Slide 3)

We use a special mathematical notation to describe complexity. This notation ignores constant factors and lower-order terms, focusing only on the *dominant term* that dictates the scaling behavior as  $N$  becomes very large.

### 1.2.1 O / Big-Oh (Worst Case Upper Bound)

This is the most common notation. It describes the **upper bound** or the **worst-case scenario**.

- **Definition:** An algorithm is  $O(f(N))$  if its runtime  $t(N)$  is *always less than* some constant  $c$  times  $f(N)$  (for a large enough  $N$ ).
- **Analogy:** Think of it as a **speed limit for slowness**. A  $O(N^2)$  algorithm is *guaranteed* to be no worse than quadratic. It might be faster (it could be  $O(N)$ ), but it will never be  $O(N^3)$ .
- **Consequences:**
  1. **We ignore constants.**  $O(2N)$  and  $O(1000N)$  are both just  $O(N)$ . Why? Because as  $N$  goes to infinity, the *linear* nature is what matters, not the slope.
  2. **We keep the dominant term.** An algorithm that takes  $N^2 + 50N + 1000$  steps is just  $O(N^2)$ .
- **Analogy for Dominant Terms:** Imagine you are building a skyscraper ( $N$  floors) and also painting the lobby. The time to build the  $N$  floors scales quadratically ( $N^2$ ). The time to paint the lobby is constant. As  $N$  (the number of floors) gets huge, the lobby-painting time becomes completely irrelevant to the total project time. The  $N^2$  term "dominates" everything else.

### 1.2.2 Ω / Big-Omega (Best Case Lower Bound)

This describes the **lower bound** or the **best-case scenario**.

- **Definition:** An algorithm is  $\Omega(f(N))$  if its runtime  $t(N)$  is *always greater than* some constant  $c$  times  $f(N)$ .
- **Analogy:** This is a **guarantee of work**. It says, "No matter how lucky you get, this algorithm will *never be faster than*  $f(N)$ ."

### 1.2.3 Θ / Big-Theta (Tight Bound)

This is used when the worst case and best case are the same.

- **Definition:** An algorithm is  $\Theta(f(N))$  if it is *both*  $O(f(N))$  and  $\Omega(f(N))$ .
- **Analogy:** This is an **exact price**. The algorithm will *always* take this amount of time to scale, no matter what the input.

### 1.3 The Real-World Impact of Complexity

(Slides 4, 5, 6)

These notations aren't just academic. They have profound, practical consequences. Let's assume a computer can perform 1 billion operations per second (1 G-op/s).

Table 1: Time taken assuming 1 G-op/s (Slide 4)

Complexity	N=10	N=100	N=1,000	N=1,000,000	N=1,000,000,000
$O(1)$	1 ns	1 ns	1 ns	1 ns	1 ns
$O(\ln N)$	3 ns	7 ns	10 ns	20 ns	30 ns
$O(N)$	10 ns	100 ns	1 $\mu$ s	1 ms	1 s
$O(N \ln N)$	30 ns	700 ns	10 $\mu$ s	20 ms	30 s
$O(N^2)$	100 ns	10 $\mu$ s	1 ms	<b>17 min</b>	<b>31.7 years</b>
$O(N^3)$	1 $\mu$ s	1 ms	1 s	<b>31.7 years</b>	$3 \times 10^{10}$ years
$O(2^N)$	1 $\mu$ s	$10^{14}$ yrs	$10^{285}$ yrs	—	—

The lesson from Table 1 is clear:

- $O(1)$ ,  $O(\ln N)$ ,  $O(N)$ , and  $O(N \ln N)$  are all exceptionally fast and scalable. We call these "efficient" algorithms.
- $O(N^2)$  (quadratic) becomes unusable very quickly. An  $N$  of one million, which is not large for modern datasets, would take 17 minutes.
- $O(N^3)$  (cubic) and  $O(2^N)$  (exponential) are "intractable." They are completely unusable for anything other than trivially small  $N$ .

**But what if we buy a faster computer?** Slide 5 shows the same table for a 10 Peta-op/s supercomputer (10,000 times faster).

- The  $O(N^3)$  algorithm for  $N = 1,000,000$  drops from 31.7 years to 100 seconds. This is a huge improvement!
- **But this is a trap.** A 10,000x faster computer only lets you solve a problem  $\sqrt[3]{10000} \approx 21.5$  times larger in the same amount of time.
- For  $N = 1,000,000,000$ , the  $O(N^3)$  algorithm *still* takes 3000 years.

**The fundamental lesson of complexity is that a good algorithm on a slow computer will always beat a bad algorithm on a fast computer for a large enough  $N$ .** You cannot fix a bad algorithm with better hardware.

When choosing (Slide 6), you *always* prefer the lower complexity class.  $O(1000 \ln N)$  is still  $O(\ln N)$ , and it will *always* be faster than  $O(N)$  once  $N$  is large enough.

## 1.4 Complexity Examples in C++

(Slides 7, 8, 9)

Let's analyze some simple C++ code blocks.

### 1.4.1 Example 1: Single Loop (Slide 7)

```
1 int n = 1000;
2 for (int i = 0; i < n; ++i) {
3     std::cout << i*i << std::endl;
4 }
```

Listing 1: A simple linear-time loop.

- The work inside the loop (multiplication, output) is constant time. We call this  $O(1)$ .
- The loop executes exactly  $n$  times.
- Total time =  $n \times O(1) = O(n)$ .
- Since the best case and worst case are identical (it *always* runs  $n$  times), we can say this is  $\Theta(n)$ .

### 1.4.2 Example 2: Nested Loop (Slide 8)

```
1 int n = 1000;
2 for (int i = 0; i < n; ++i) {
3     for (int j = 0; j < i; ++j) {
4         std::cout << i*j << std::endl;
5     }
6 }
```

Listing 2: A nested quadratic-time loop.

- The outer loop runs  $n$  times (for  $i = 0, 1, \dots, n - 1$ ).
- The inner loop runs  $i$  times.
- The total number of operations is the sum of  $0 + 1 + 2 + 3 + \dots + (n - 1)$ .
- This is a famous arithmetic series, which sums to  $\frac{n(n-1)}{2} = \frac{n^2-n}{2}$ .
- Using our notation rules, we drop the constants ( $\frac{1}{2}$ ) and the lower-order term ( $-n/2$ ).
- The complexity is  $O(n^2)$ .
- Since it always does this, it is  $\Theta(n^2)$ .

### 1.4.3 Example 3: Search with 'break' (Slide 9)

```
1 // Block 1: Allocation
2 int n = 1000;
3 double* x = new double[n]; // O(n) operation
4 for (int i=0; i < n; ++i) // O(n) operation
5     x[i] = i;
6
7 // Block 2: Search
8 int pos = -1;
9 for (int i=0; i < n; ++i) {
10     if (x[i] == y) {
11         pos = i;
12         break; // This is the important line!
13     }
14 }
```

Listing 3: A linear search with different cases.

- **Block 1:** Allocating  $n$  elements is  $O(n)$ . The loop to fill them is  $O(n)$ . Total time is  $O(n) + O(n)$ , which is still just  $O(n)$ .
- **Block 2:** The 'break' statement changes everything.
  - **Worst Case:** The value 'y' is the last element ( $x[n - 1]$ ) or is not in the array at all. The loop runs  $n$  times. The complexity is  $O(n)$ .
  - **Best Case:** The value 'y' is the very first element ( $x[0]$ ). The loop runs once, hits the 'break', and exits. The complexity is  $O(1)$ .
- **Conclusion:** Because the best and worst cases are different, we cannot use  $\Theta$ . We say the algorithm has a worst-case of  $O(n)$  and a best-case of  $\Omega(1)$ .

## 1.5 Amortized Analysis: The "Clever Way" to Grow Arrays

(Slides 10, 11)

**The "Simple Way" (Slide 10)** Imagine you have a fixed-size array of  $N$  elements, and you want to add one more. You must:

1. Allocate a *new* array of size  $N + 1$ .
2. *Copy* all  $N$  elements from the old array to the new one.
3. Add the new  $(N + 1)^{th}$  element.

The copy step takes  $O(N)$  time. If you do this for *every single element* you add, adding  $N$  elements one by one will take  $O(1) + O(2) + \dots + O(N) = O(N^2)$  time. This is horribly slow.

**The "Clever Way" (Amortized  $O(1)$ ) (Slide 11)** This is the strategy used by C++'s 'std::vector'.

1. When you ask for an array, it secretly allocates *extra space*. This is its **capacity**. The number of elements you are using is its **size**.
2. As you add elements, it just increments the 'size'. This is a  $O(1)$  operation.
3. When you try to add an element and 'size == capacity', the vector is full. It now performs an "expensive" operation:
4. It allocates a *new*, much larger array, typically **double the size** ( $2N$ ).
5. It copies the  $N$  old elements. This one operation is  $O(N)$ .
6. It adds the new element.

This single  $O(N)$  operation seems bad, but it just bought you  $N$  more "cheap"  $O(1)$  additions.

**Amortized Analogy:** You have a 10-person dining table.

- **Simple Way ( $O(N)$ ):** When guest #11 arrives, you move everyone to an 11-person table. When guest #12 arrives, you move everyone to a 12-person table. This is a nightmare.
- **Clever Way (Amortized  $O(1)$ ):** When guest #11 arrives, you move all 10 people to a 20-person table. This one move is expensive ( $O(N)$ ). But now, guests #12 through #20 can just sit down instantly ( $O(1)$ ).

The total cost of  $N$  additions is  $O(N)$  (for the one expensive copy) plus  $N \times O(1)$  (for the cheap additions). The total cost is  $O(N)$ . The *average* or *amortized* cost per operation is  $O(N)/N = \mathbf{O(1)}$ .

## 2 The C++ Standard Template Library (STL)

### 2.1 What is the STL?

(Slides 12, 13)

The **Standard Template Library (STL)** is a powerful, efficient, and well-tested library of code built into C++. Its key idea is **generic programming**: writing code that works with any data type.

The STL is built on three core pillars:

1. **Containers:** Data structures that store your data.
2. **Algorithms:** Functions that process your data.
3. **Iterators:** The "glue" that connects algorithms to containers.

**Analogy:**

- **Containers** are your filing cabinets, phone books, and to-do lists.

- **Algorithms** are the *actions* you take: ‘sort()’ the files, ‘find()’ a phone number, ‘reverse()’ the to-do list.
- **Iterators** are a “generic hand” that knows how to point to an item and move to the next one, regardless of whether it’s in a filing cabinet or a phone book.

This design is brilliant: the ‘std::sort’ algorithm doesn’t know what a ‘std::vector’ is. It only knows how to use iterators. This means ‘std::sort’ was written *once* and works on almost any container.

## 2.2 Simple Utilities

(Slides 14, 15)

**‘std::string’ (Slide 14)** Found in ‘<string>’, this is the C++ class for handling text. ‘std::wstring’ is for “wide” characters (e.g., for non-English alphabets). It’s a container-like class with many useful member functions.

**‘std::pair’ (Slide 15)** Found in ‘<utility>’, ‘std::pair’ is a simple template class that just holds two values, which can be of different types.

```

1 #include <utility>
2 #include <string>
3 #include <iostream>
4
5 int main() {
6     // Create a pair of a string and an integer
7     std::pair<std::string, int> student("Alice", 20);
8
9     // Access the elements using .first and .second
10    std::cout << "Name: " << student.first << std::endl;
11    std::cout << "Age: " << student.second << std::endl;
12
13    // You can also use std::make_pair
14    auto student2 = std::make_pair("Bob", 22);
15 }
```

Listing 4: Using std::pair

This is very useful for functions that need to return two values, or as the element type for ‘std::map’. For more than two items, C++ offers ‘std::tuple’.

## 3 STL Container Deep Dive

(Slide 16)

The STL provides many containers, which we can group into categories. We will explore the most important ones.

### 3.1 Common Container Operations

All STL containers share a common interface for basic operations:

```
1 // Works for vector, list, deque, set, map, etc.
2 std::vector<int> container;
3
4 // Size and capacity
5 size_t sz = container.size();           // Number of elements
6 bool empty = container.empty();         // Returns true if size() == 0
7 size_t max_sz = container.max_size();   // Maximum possible size
8
9 // Iterators (every container provides these!)
10 auto it_begin = container.begin();     // Iterator to first element
11 auto it_end = container.end();          // Iterator past last element
12 auto cit_begin = container.cbegin();   // Const iterator to first
13 auto cit_end = container.cend();       // Const iterator past last
14
15 // Clear all elements
16 container.clear(); // Removes all elements, size() becomes 0
17
18 // Swap contents with another container (very fast - just swaps
19 // pointers)
20 std::vector<int> other;
21 container.swap(other); // Or: std::swap(container, other);
```

Listing 5: Universal Container Operations

### 3.2 Sequence Containers

These containers store elements in a specific linear order that you define.

#### 3.2.1 ‘std::vector’ — The “Smart Array”

(Slides 17, 18, 19, 44, 47)

This is the most common and useful container. It is the C++ “smart array” and should be your default choice. Include with ‘<vector>’.

- **Internal Structure:** A single, contiguous block of memory (Slide 17).
- **Pros:**
  - **$O(1)$  Random Access:**  $v[i]$  is just a pointer calculation.
  - **Excellent Cache Locality:** Because all data is side-by-side, the CPU can pre-load elements into its ultra-fast cache, making loops very fast.
  - **Amortized  $O(1)$  push\_back():** Adding to the end is (on average)  $O(1)$  (Slide 19).
- **Cons:**
  - **$O(N)$  Insert/Erase:** Adding or removing an element in the *middle* is very slow (Slide 18). You have to “shift” all elements after that point, which is an  $O(N)$  copy operation.

```

1 #include <vector>
2
3 // === Construction ===
4 std::vector<int> v1;                                // Empty vector
5 std::vector<int> v2(5);                            // 5 elements, default value
6 (0)
7 std::vector<int> v3(5, 42);                         // 5 elements, all set to 42
8 std::vector<int> v4 = {10, 20, 30};                // Initialize with list
9 std::vector<int> v5(v4);                           // Copy constructor
9 std::vector<int> v6(v4.begin(), v4.begin() + 2); // From range [10,20]

```

Listing 6: std::vector Construction

```

1 std::vector<int> v = {10, 20, 30};
2
3 // === Adding Elements ===
4 v.push_back(40);           // Add to end: [10,20,30,40] - O(1) amortized
5 v.insert(v.begin(), 5);   // Insert at beginning: [5,10,20,30,40] - O(N)
6 v.insert(v.begin() + 2, 15); // Insert at position 2: [5,10,15,20,30,40]
- O(N)
7 v.emplace_back(50);       // Construct in-place (more efficient) - O(1)
amortized
8
9 // === Removing Elements ===
10 v.pop_back();            // Remove last: [5,10,15,20,30] - O(1)
11 v.erase(v.begin());      // Remove first: [10,15,20,30] - O(N)
12 v.erase(v.begin() + 1);  // Remove at position 1: [10,20,30] - O(N)
13 v.erase(v.begin(), v.begin() + 2); // Remove range: [30] - O(N)
14 v.clear();               // Remove all elements - O(N)

```

Listing 7: std::vector Adding and Removing Elements

```

1 std::vector<int> v = {10, 20, 30, 40};
2
3 int first = v.front();        // 10 - First element
4 int last = v.back();         // 40 - Last element
5 int val = v[1];              // 20 - FAST but UNSAFE (no bounds check)
6 int val_safe = v.at(1);      // 20 - SAFE (throws exception if out of
bounds)
7
8 // Using iterators
9 int first_it = *v.begin();  // 10 - First element via iterator
10 int last_it = *(v.end() - 1); // 40 - Last element via iterator

```

Listing 8: std::vector Element Access

```

1 std::vector<int> v = {10, 20, 30, 40};
2
3 size_t s = v.size();          // 4 - Number of elements currently stored
4 size_t c = v.capacity();      // >= 4 - Space allocated (may be larger)
5 bool empty = v.empty();       // false - Equivalent to size() == 0
6
7 // Pre-allocate space (optimization!)
8 v.reserve(100); // Guarantee capacity >= 100 (no reallocation until
then)
9 v.shrink_to_fit(); // Request to reduce capacity to fit size
10
11 // Change size

```

```

12 v.resize(6);           // Size becomes 6, new elements are 0:
13   [10,20,30,40,0,0]
14 v.resize(3);           // Size becomes 3: [10,20,30] (elements removed)
15 v.resize(5, 99);       // Size becomes 5, new elements are 99:
16   [10,20,30,99,99]

```

Listing 9: std::vector Size and Capacity Management

```

1 std::vector<int> v = {10, 20, 30, 40, 50};
2
3 // Get iterators
4 auto it_begin = v.begin();      // Points to first element (10)
5 auto it_end = v.end();          // Points PAST last element (invalid to
6   dereference!)
7 auto rit_begin = v.rbegin();    // Reverse: points to last (50)
8 auto rit_end = v.rend();        // Reverse: points before first
9
10 // Iterate and modify
11 for (auto it = v.begin(); it != v.end(); ++it) {
12     *it *= 2; // Double each element
13 }
14 // v is now [20, 40, 60, 80, 100]
15
16 // Const iteration (read-only)
17 for (auto cit = v.cbegin(); cit != v.cend(); ++cit) {
18     std::cout << *cit << " ";
19     // *cit = 0; // ERROR: Cannot modify through const iterator
20 }
21
22 // Modern range-based for loop (preferred when you don't need iterator
23 // itself)
24 for (auto& elem : v) {         // Reference to modify
25     elem += 10;
26 }
27 for (const auto& elem : v) { // Const reference for read-only
28     std::cout << elem << " ";
29 }

```

Listing 10: std::vector Iterator Operations

### 3.2.2 std::deque — The "Double-Ended Queue"

(Slide 20)

deque (pronounced "deck") is very similar to vector but with one extra power: it supports fast insertion/removal at *both the front and the back*. Include with <deque>.

- **Key Feature:**  $O(1)$  push\_front(), pop\_front(), push\_back(), pop\_back().
- **Analogy:** A line of people where you can add/remove from both the front and the back of the line instantly.

- **Internals:** More complex than `vector` (often a list of small arrays). It still provides  $O(1)$  random access `d[i]`, but it's slightly slower than `vector`'s access due to bad cache locality.
- **Use Case:** Use `deque` if you *know* you need to add/remove from the front. Otherwise, `vector` is usually faster.

### 3.2.3 `std::list` — The "Linked List"

(Slides 24, 49)

`list` is a **doubly-linked list**. It is fundamentally different from `vector`. Include with `<list>`.

- **Internal Structure:** A chain of "nodes." Each node contains a value and two pointers: one to the 'next' node and one to the 'previous' node.
- **Pros:**
  - **$O(1)$  Insert/Erase:** If you have an iterator pointing to a location, you can insert or remove an element in  $O(1)$  time. You just re-wire the pointers of the neighbors. No element shifting is needed.
- **Cons:**
  - **$O(N)$  Access:** You *cannot* do `l[i]`. To get the 100th element, you *must* start at the beginning and follow the `next` pointer 100 times.
  - **Bad Cache Locality:** The nodes can be scattered all over memory, which is slow for the CPU's cache.

**Analogy:**

- `vector`: A row of mailboxes. Fast to jump to mailbox #100 ( $O(1)$ ), but slow to add a new mailbox in the middle ( $O(N)$ ).
- `list`: A treasure hunt. Fast to add a new clue in the middle ( $O(1)$ ), but to find clue #100, you must follow the first 99 clues ( $O(N)$ ).

Because `list` is so different, it has its own special member functions (Slide 49) that are much faster than the generic algorithms:

- `l.sort()`: A special  $O(N \ln N)$  sort that just re-wires pointers, never copying values.
- `l.splice(it, other_list)`: Moves all nodes from `other_list` into `l` at position `it`. This is a  $O(1)$  operation!
- `l.remove(value)`: Removes all elements equal to `value`.
- `l.remove_if(predicate)`: Removes all elements for which the predicate function returns `true`. (This is key for the Penna model!)

### 3.3 Container Adapters

(Slides 21, 22, 23, 48)

Adapters are not new containers. They are "wrappers" that provide a simpler, more restrictive interface on top of an existing container (usually 'std::deque' by default).

**Analogy:** A Pez dispenser. It's just a wrapper around a stack of candy, but it *restricts* you. You can only 'push()' (load) from one end and 'pop()' (eat) from the other. This restriction is a *feature*, as it enforces a specific access pattern.

#### 3.3.1 std::stack

(Slide 21) Implements a **LIFO (Last-In, First-Out)** structure. Include with <stack>.

- **Analogy:** A stack of plates. You put a plate on top, you take a plate from the top.
- **Operations ( $O(1)$ ):**
  - `s.push(value)`: Add an element to the top.
  - `s.pop()`: Remove the top element.
  - `s.top()`: Get a reference to the top element.

#### 3.3.2 std::queue

(Slide 22) Implements a **FIFO (First-In, First-Out)** structure. Include with <queue>.

- **Analogy:** A checkout line (or "queue") at a store.
- **Operations ( $O(1)$ ):**
  - `q.push(value)`: Add an element to the *back*.
  - `q.pop()`: Remove the element from the *front*.
  - `q.front()`: Get a reference to the front element.
  - `q.back()`: Get a reference to the back element.

#### 3.3.3 std::priority\_queue

(Slide 23) A special queue where elements are removed based on priority, not arrival time. Include with <queue>.

- **Analogy:** An Emergency Room waiting line. Patients are seen by severity (priority), not by who arrived first.
- By default, "priority" means "largest value."
- **Operations:**
  - `pq.push(value)`: Adds an element, sorting it into the queue. ( $O(\log N)$ )
  - `pq.pop()`: Removes the *highest priority* element. ( $O(\log N)$ )
  - `pq.top()`: Get a reference to the highest priority element. ( $O(1)$ )

## 3.4 Associative Containers (Trees)

(Slides 25, 26, 27, 28, 50)

We have a problem:

- `vector`: Fast access ( $O(1)$ ), but slow search ( $O(N)$ ) in an unsorted vector).
- `list`: Fast insert/erase ( $O(1)$ ), but slow access and search ( $O(N)$ ).

What if we need **fast search, fast insert, AND fast erase?**

**Solution: A Balanced Binary Search Tree (BST).**

- **Structure (Slide 26, 27):** A tree made of nodes. Each node has a value, a pointer to a 'left' child (with a smaller value) and a 'right' child (with a larger value).
- **Performance:** The tree is automatically "balanced" to keep it bushy, not stringy. This guarantees that the height of the tree is  $\log N$ .
- This means **search, insert, and erase are all  $O(\log N)$ .**
- $O(\log N)$  (logarithmic) is *extremely* fast (see Table 1).

**Analogy:** The game of "20 Questions." You start at the root and ask a "smaller or larger?" question at each node, dividing the remaining search space in half each time.

The STL provides two main tree-based containers:

### 3.4.1 `std::set`

Include with `<set>`.

- **What it is:** Stores a collection of **unique, sorted** keys.
- **Analogy:** A VIP guest list. It's sorted alphabetically, and you can't be on the list twice.

- **Use Case:** When you just need to know if an item *exists* in a set, and you need to do it quickly.

- **Syntax:**

```

1 #include <set>
2 std::set<std::string> banned_users;
3 banned_users.insert("Alice");
4 banned_users.insert("Bob");
5 banned_users.insert("Alice"); // This does nothing, "Alice" is
   already in.
6
7 // Fast O(log N) lookup
8 if (banned_users.count("Bob") > 0) {
9     // ...
10}
11

```

Listing 11: Using std::set

### 3.4.2 std::map

Include with `<map>`.

- **What it is:** Stores a collection of **unique, sorted key-value pairs**.
- **Analogy:** A dictionary or a phone book. The "key" is the word (e.g., "algorithm"), and the "value" is the definition. The keys are sorted.
- **Use Case:** Associating one piece of data with another.

- **Syntax:**

```

1 #include <map>
2 #include <string>
3
4 std::map<std::string, int> student_ages;
5
6 // Insert using [] operator (O(log N))
7 student_ages["Charlie"] = 21;
8 student_ages["David"] = 19;
9
10 // Insert using .insert()
11 student_ages.insert(std::make_pair("Eve", 23));
12
13 // Fast O(log N) lookup
14 std::cout << "David's age: " << student_ages["David"] << std::
   endl;
15

```

Listing 12: Using std::map

**Note:** multiset and multimap also exist, which allow duplicate keys.

## 4 The Magic Glue: Iterators

### 4.1 The $N \times M$ Problem

(Slides 30-34)

We have a problem. We have  $M$  containers (`vector`, `list`, `deque`...) and  $N$  algorithms (`find`, `copy`, `sort`...).

- To loop through a `vector`, you use a pointer: `for (T* p = ...)`
- To loop through a `list`, you use a node: `for (node* p = ...)`

The code is different! (Slide 30). Does this mean we have to write  $N \times M$  different functions (e.g., `find_in_vector`, `find_in_list`, `sort_vector`, `sort_list`)? This would be a nightmare (Slides 31-34).

### 4.2 The Solution: Generic Traversal

(Slides 35, 36)

The answer is **NO**. The STL solves this with **Iterators**. An iterator is an object that acts like a "generic pointer." It abstracts away the details of the container.

Every container provides two functions:

- `container.begin()`: Returns an iterator to the *first* element.
- `container.end()`: Returns an iterator *past the last* element.

All iterators, no matter what container they come from, support common operations:

- `++it`: Move to the next element.
- `*it`: Get the value of the element (*dereference*).
- `it1 == it2`: Compare two iterators.

Now, we can write **one** generic loop that works on **any** container:

```
1 // This code works if 'c' is a vector, a list, a deque, or a set!
2 for (auto it = c.begin(); it != c.end(); ++it) {
3     auto value = *it;
4     // ... do something with value ...
5 }
```

Listing 13: The generic iterator loop (Slide 35)

This is so common that C++11 introduced a "range-based for loop" that is just syntactic sugar for the code above:

```
1 // This is the preferred, modern way to loop
2 for (auto const& element : c) {
3     // ... do something with element ...
4 }
```

Listing 14: The modern C++11 loop (Slide 36)

## 4.3 How Iterators are Implemented

(Slides 37, 38)

This "generic pointer" is just an abstraction.

- **For `std::vector` (Slide 37):** An iterator is just a raw pointer. `begin()` returns a `T*` to the first element. `++it` is just pointer arithmetic.
- **For `std::list` (Slide 38):** An iterator is a small class that holds a pointer to a node. This class **overloads the operators** to pretend to be a pointer:
  - Its `operator++()` function is defined to mean `p = p->next`.
  - Its `operator*()` function is defined to mean `return p->value`.

The algorithm doesn't know or care about this difference. It just calls `++it` and `*it`, and the magic works.

## 4.4 Complete Iterator Operations Reference

### 4.4.1 Essential Iterator Operators

```
1 std::vector<int> v = {10, 20, 30, 40};  
2  
3 // Get iterator to first element  
4 auto it = v.begin();  
5  
6 // Get iterator past the last element (DO NOT dereference!)  
7 auto it_end = v.end();  
8  
9 // Const iterators (read-only)  
10 auto cit = v.cbegin(); // Points to first element (const)  
11 auto cit_end = v.cend(); // Points past last element (const)  
12  
13 // Reverse iterators (iterate backwards)  
14 auto rit = v.rbegin(); // Points to last element  
15 auto rit_end = v.rend(); // Points before first element
```

Listing 15: Core Iterator Operations: Obtaining Iterators

```
1 std::vector<int> v = {10, 20, 30};  
2 auto it = v.begin();  
3  
4 // Dereference to get value  
5 int value = *it; // value = 10  
6  
7 // Access member of pointed object (for objects/structs)  
8 std::vector<std::string> names = {"Alice", "Bob"};  
9 auto name_it = names.begin();  
10 int length = name_it->length(); // Calls string::length()  
11 // Equivalent to: (*name_it).length()
```

Listing 16: Dereferencing - Accessing Values

```

1 std::vector<int> v = {10, 20, 30, 40};
2 auto it = v.begin();
3
4 // Pre-increment (preferred for iterators)
5 ++it; // Now points to 20
6 int val1 = *it; // val1 = 20
7
8 // Post-increment (creates a copy, slightly less efficient)
9 auto old_it = it++; // it moves to 30, old_it still at 20
10 int val2 = *old_it; // val2 = 20
11 int val3 = *it; // val3 = 30
12
13 // Pre-decrement (for bidirectional iterators)
14 --it; // Back to 20
15
16 // Post-decrement
17 it--; // Back to 10

```

Listing 17: Incrementing and Decrementing Iterators

```

1 std::vector<int> v = {10, 20, 30};
2 auto it1 = v.begin();
3 auto it2 = v.begin();
4 auto it_end = v.end();
5
6 // Equality comparison
7 if (it1 == it2) {
8     // True: both point to same position
9 }
10
11 // Inequality comparison (most common for loops!)
12 if (it1 != it_end) {
13     // True: it1 is not past-the-end
14     int value = *it1; // Safe to dereference
15 }
16
17 // IMPORTANT: Always check iterator != end() before dereferencing!
18 while (it1 != v.end()) {
19     std::cout << *it1 << " ";
20     ++it1;
21 }

```

Listing 18: Comparison Operators - Essential for Loops!

```

1 std::vector<int> v = {10, 20, 30, 40, 50};
2 auto it = v.begin();
3
4 // Addition (jump forward)
5 auto it_plus_3 = it + 3; // Points to 40
6 int val = *it_plus_3; // val = 40
7
8 // Subtraction (jump backward)
9 auto it_minus_1 = it_plus_3 - 1; // Points to 30
10
11 // Compound assignment
12 it += 2; // Move forward 2 positions (now at 30)
13 it -= 1; // Move backward 1 position (now at 20)

```

```

14 // Array-like indexing
15 int value = it[2]; // Access element 2 positions ahead
16 // Equivalent to *(it + 2)
17
18 // Distance between iterators
19 auto dist = v.end() - v.begin(); // dist = 5 (size of vector)
20
21 // Relational comparisons
22 auto it1 = v.begin();
23 auto it2 = v.begin() + 2;
24 if (it1 < it2) { // True: it1 comes before it2
25     // ...
26 }
27 }
```

Listing 19: Random Access Iterator Operations (vector, deque only)

#### 4.4.2 Complete Example: Using All Iterator Operations

```

1 #include <vector>
2 #include <iostream>
3
4 int main() {
5     std::vector<int> numbers = {10, 20, 30, 40, 50};
6
7     // === Basic Iteration ===
8     std::cout << "Forward iteration:\n";
9     for (auto it = numbers.begin(); it != numbers.end(); ++it) {
10         std::cout << *it << " "; // Dereference to get value
11     }
12     std::cout << "\n";
13
14     // === Reverse Iteration ===
15     std::cout << "Reverse iteration:\n";
16     for (auto rit = numbers.rbegin(); rit != numbers.rend(); ++rit) {
17         std::cout << *rit << " ";
18     }
19     std::cout << "\n";
20
21     // === Random Access ===
22     auto it = numbers.begin();
23     std::cout << "Element at begin: " << *it << "\n"; // 10
24     std::cout << "Element at begin+2: " << *(it + 2) << "\n"; // 30
25     std::cout << "Element using []: " << it[3] << "\n"; // 40
26
27     // === Modifying Through Iterator ===
28     auto modify_it = numbers.begin();
29     *modify_it = 15; // Changes first element from 10 to 15
30
31     // === Iterator Arithmetic ===
32     auto start = numbers.begin();
33     auto end = numbers.end();
34     auto distance = end - start; // Number of elements: 5
35 }
```

```

36     auto middle = start + (distance / 2); // Points to middle element
37     std::cout << "Middle element: " << *middle << "\n"; // 30
38
39     // === Const Iterators (Read-Only) ===
40     for (auto cit = numbers.cbegin(); cit != numbers.cend(); ++cit) {
41         std::cout << *cit << " "; // Can read
42         // *cit = 100; // ERROR: Cannot modify through const iterator
43     }
44     std::cout << "\n";
45
46     return 0;
47 }
```

Listing 20: Comprehensive iterator usage example

## 4.5 Iterator Categories

(Slide 39)

Not all iterators are created equal. They are categorized by their "power."

### 4.5.1 Iterator Category Hierarchy

1. **Input/Output Iterator:** Weakest. Can only move forward and be read-/written once.
  - **Supported:** `++it, it++, *it, ==, !=`
  - **Example:** `std::istream_iterator`
2. **Forward Iterator:** Can move forward (`++`) many times.
  - **Supported:** All Input operations, plus multi-pass guarantee
  - **Example:** `std::forward_list::iterator`
  - **Use case:** `std::find()`, `std::replace()`
3. **Bidirectional Iterator:** Can move forward (`++`) and backward (`-`).
  - **Supported:** All Forward operations, plus `-it, it-`
  - **Examples:** `std::list::iterator`, `std::set::iterator`, `std::map::iterator`
  - **Use case:** `std::reverse()`, `std::find_end()`
4. **Random Access Iterator:** Most powerful. Can jump to any position in  $O(1)$  time.
  - **Supported:** All Bidirectional operations, plus:
    - `it + n, it - n, it += n, it -= n`
    - `it[n]` (array subscript)

- `it2 - it1` (distance)
- `<, >, <=, >=` (comparison)
- **Examples:** `std::vector::iterator`, `std::deque::iterator`, raw pointers
- **Use case:** `std::sort()`, `std::binary_search()`

```

1 // Forward Iterator (std::forward_list)
2 std::forward_list<int> fl = {1, 2, 3, 4, 5};
3 auto flit = fl.begin();
4 ++flit;           // OK: Can move forward
5 // --flit;         // ERROR: Cannot move backward
6 // flit + 2;      // ERROR: Cannot jump
7
8 // Bidirectional Iterator (std::list)
9 std::list<int> lst = {1, 2, 3, 4, 5};
10 auto lit = lst.begin();
11 ++lit;           // OK: Can move forward
12 --lit;           // OK: Can move backward
13 // lit + 2;      // ERROR: Cannot jump
14
15 // Random Access Iterator (std::vector)
16 std::vector<int> vec = {1, 2, 3, 4, 5};
17 auto vit = vec.begin();
18 ++vit;           // OK: Can move forward
19 --vit;           // OK: Can move backward
20 vit = vit + 3;  // OK: Can jump
21 int val = vit[1]; // OK: Can use subscript

```

Listing 21: Iterator Category Example

**Important:** Algorithms specify the *minimum* category they need. `std::find` only needs a Forward iterator. `std::sort` requires a Random Access iterator (which is why you can't call `std::sort` on a `std::list`).

## 5 The Generic Algorithms

(Slides 53-63)

Now we get the payoff. The `<algorithm>` header contains dozens of pre-built, highly-optimized functions that operate on iterators. **You should always prefer these to writing your own loops.**

### 5.1 Example: `std::find`

(Slide 54) `std::find` searches a range for a value.

```

1 #include <vector>
2 #include <algorithm>
3 #include <iostream>
4
5 std::vector<int> v = {10, 20, 30, 40};

```

```

6 int value_to_find = 30;
7
8 // find returns an iterator
9 auto it = std::find(v.begin(), v.end(), value_to_find);
10
11 // --- CRITICAL CHECK ---
12 // If not found, find returns the .end() iterator!
13 if (it != v.end()) {
14     std::cout << "Found it! Value is " << *it << std::endl;
15 } else {
16     std::cout << "Value not found." << std::endl;
17 }

```

Listing 22: Using std::find

The implementation of ‘find’ is just the simple generic loop from Slide 35.

## 5.2 Example: std::find\_if

(Slide 55) `find_if` is more powerful. Instead of a value, it takes a **predicate**: a function (or function-like object) that returns `bool`.

```

1 // A predicate function
2 bool isEven(int x) {
3     return x % 2 == 0;
4 }
5
6 std::vector<int> v = {1, 3, 5, 6, 7, 9};
7
8 // Find the first element for which isEven() returns true
9 auto it = std::find_if(v.begin(), v.end(), isEven);
10
11 if (it != v.end()) {
12     std::cout << "First even number is " << *it << std::endl; // Prints 6
13 }

```

Listing 23: Using std::find<sub>i</sub>f

## 5.3 Tricky Case 1: Member Functions as Predicates

(Slide 56)

What if your predicate is a *member function* of a class?

```

1 class Animal {
2 public:
3     bool is_pregnant() const;
4     // ...
5 };
6
7 std::list<Animal> flock;
8 // ...
9 // This will NOT compile!

```

```

10 auto it = std::find_if(flock.begin(), flock.end(), &Animal::
    is_pregnant);

```

The `find_if` algorithm doesn't know how to call a member function. It expects a global function `bool(Animal)`.

**Solution:** Use an adapter from `<functional>` called `std::mem_fn`.

```

1 #include <functional> // Need this!
2
3 // This works!
4 auto it = std::find_if(flock.begin(), flock.end(),
5                         std::mem_fn(&Animal::is_pregnant));

```

## 5.4 Tricky Case 2: Copying into Empty Containers

(Slide 57)

`std::copy` copies elements from one range to a destination iterator.

```

1 std::vector<int> v = {1, 2, 3};
2 std::vector<int> w; // w is EMPTY!
3
4 // This will CRASH!
5 // w.begin() points to nothing, so copy tries to write to
6 // invalid memory.
7 std::copy(v.begin(), v.end(), w.begin());

```

**Solution 1 (Clumsy):** `w.resize(v.size());` first.

**Solution 2 (Elegant):** Use a `std::back_inserter` from `<iterator>`.

```

1 #include <iterator> // Need this!
2
3 // This works!
4 // std::back_inserter(w) is an iterator adapter that
5 // "pretends" to be a normal iterator, but its "write"
6 // operation (operator=) actually calls w.push_back().
7 std::copy(v.begin(), v.end(), std::back_inserter(w));
8
9 // w is now {1, 2, 3}

```

## 5.5 Algorithm Naming Conventions

(Slide 62)

The algorithm names are very consistent.

- Suffix `_if`: Takes a predicate instead of a value.
  - `find(..., val)`
  - `find_if(..., pred)`
- Suffix `_copy`: Does not modify the original range. Writes a *copy* of the result to a destination.

- `reverse(beg, end)`: Reverses the range in-place.
- `reverse_copy(beg, end, dest)`: Writes a reversed copy to `dest`, leaving `[beg, end]` unchanged.
- Suffix `_copy_if`: Combines both.
  - `remove_copy_if(beg, end, dest, pred)`: Copies all elements *except* those for which `pred` is true.

For a full list of all algorithms, see a C++ reference like ‘[cppreference.com](#)’.

## 6 Application and Summary

### 6.1 Application: The Penna Model

(Slides 58, 64)

The exercise for this week is to code the `Population` class for the Penna model. You can (and should) use the STL to make this trivial.

- The `Population` can just be a `std::list<Animal>`. We use `list` because we expect to remove many animals from the middle (when they die), and `list` is  $O(1)$  for this.
- To remove all dead animals, you don’t need to write your own loop. You can use the `list`’s special `remove_if` function combined with the `mem_fn` adapter.

```

1 #include <list>
2 #include <functional>
3 #include "animal.hpp"
4
5 class Population {
6 public:
7     void remove_dead() {
8         // This one line replaces an entire, complex,
9         // error-prone for-loop.
10        pop_.remove_if(std::mem_fn(&Animal::is_dead));
11    }
12    // ...
13 private:
14     std::list<Animal> pop_;
15 };

```

The challenge is to write the entire `Population` class **without any raw loops** (`for`, `while`), using only STL algorithms. This greatly increases reliability.

## 6.2 Summary

(Slide 65)

- **Rule 1:** Before you write any code, **check the C++ standard library.** `find`, `sort`, `vector`, `map`, etc., are already written, heavily optimized, and bug-free.
- **Rule 2:** When you design your own classes, try to emulate the STL. Provide `.begin()` and `.end()` iterators so your classes can be used with generic algorithms.
- **Rule 3:** Don't be scared by the long error messages. Template metaprogramming (the `<...>` syntax) can produce huge, unreadable errors. This is normal. Look at the *first* line of the error; that's usually where the real problem is.

## Part II

# Plotting / Scientific Visualization (week07b)

## 7 Plotting Our Data

(Slide 2)

Our scientific simulations will produce large amounts of data (e.g., population size over time). A text file full of numbers is useless for understanding; we need to *visualize* it. This section introduces the tools we can use.

### 7.1 Gnuplot

- **Website:** <http://www.gnuplot.info/>
- **What it is:** A very old, powerful, and stable command-line plotting program.
- **Pros:** It's fast, universal (Linux, macOS, Windows), and excellent for generating 2D and 3D plots quickly, often from within a script.

### 7.2 Python + Matplotlib

- **Website:** <https://www.python.org/> and <https://matplotlib.org/>
- **What it is:** The combination of the Python programming language and its most popular plotting library, Matplotlib.
- **Pros:** This is the *de facto* standard in many scientific fields. It is extremely flexible, powerful, and can create publication-quality graphs.

### 7.3 Best Practice: Python Virtual Environments

When using Python, it is highly recommended to use **virtual environments** (e.g., via Python's built-in 'venv' module).

- **What it is:** A tool that creates an isolated, self-contained "bubble" for each of your projects.
- **Analogy:** A virtual environment is like a separate, clean workshop for each project. For "Project A," you can install a special "bandsaw" (e.g., 'matplotlib version 3.0'). For "Project B," you can install a different "bandsaw" (e.g., 'matplotlib version 3.5'). These don't interfere. This prevents a library update for one project from breaking all your other projects.
- **How to use (basic):**

```
# 1. Create a virtual environment named "my-project-env"  
python3 -m venv my-project-env  
  
# 2. Activate it (on Linux/macOS)  
source my-project-env/bin/activate  
  
# 3. Install libraries. They will only be installed inside this "bubble"  
pip install matplotlib  
pip install numpy  
  
# 4. Run your script  
python my_simulation.py  
  
# 5. Deactivate when you are done  
deactivate
```

The lecture will cover Python in more detail later. For now, demo code can be found in the ‘week07/plotting’ directory. For 3D visualization of complex simulation data, we will later look at powerful tools like **ParaView** and **VisIt**.